

16-9-2020

FIRI Minal

FAISAL Ahsan

## PDC

- Distributed = largest data      All processor
- CC = Cache Coherence = Auto update ↑ if one processor  
Non Unified = separate data memory
- Unified Memory = shared memory

MPI = Bewerke Cluster own VMs.

↳ Distributed Memory.

→ Hybrid Distributed-Shared Memory.

Problems:

Word Count 10TB

Data Find → ↗

Draw Backs:

Data Duplication

Semaphores: used to handle race condition.  
↳ Pinning Table.

Tasks:

+ , - , / , \*

Thread Model:

Hav task ko aag aag processor pr divide kr lgs.

GPU:

(blocks, threads)

↓

64 × 64

1 block has 64 threads

CPU → Parallel Processor.  
GPU → Massive Parallel Processor.

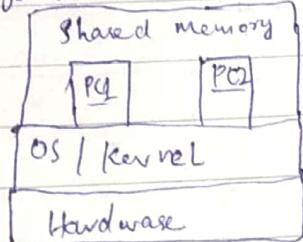
→ Thread Model Implementations:  
Pthread & openMP

→ MPI: Message Passing.

Send();

Receive();

Hybrid:



→ Data Parallel Model: (One task i.e) counting

Any Dataset 100 GB

→ 19 \* tasks on data.

OpenMP = Hybrid Model.

⇒ Task Parallelism:

Multiple task on data.

# Parallel Programming Platforms

## CH #2

Throughput: data process per unit time.

{ Data Intensive: youtube (video processing). } Scope of parallelism.  
Server Applications: video games.  
Scientific Applications: Image processing, Biometric verification.

Types OF Instructions Executions:

- ① Sequential Instruction.
- ② Parallel + Addition
- ③ Pipeline " Matrix Multiplication
- ④ Jump (switch case)
- ⑤ Branch (If - else)
- ⑥ Memory Prediction

functional units = processors.

SUPERSCALER EXECUTION: Fetch decode execute

No functional units utilized = vertical waste

Some = horizontal waste.

VLIW = Very Long Instructions word = instructions are

processors. packed & dispatched together = parallel concurrent

↳ limited to 4-way & 8-way parallel execution.

Inorder:  $M \times M$  1st row 1st column

Dynamic: Packets, file sending  
chunk order, out of order

To

$$B^x = 3 \quad A' = 3 \quad A = 3.67 \\ B^t = 4.$$

limitations of Memory System Performance.  
→ largely captured by two parameters.

latency = delay , Bandwidth = memory access per unit time.

Effective Memory Latency Time.  
→ Using TLB, cache.

→ Multithread  
    ↳ divide program in multiply threads.

→ Prefetching for latency hiding:  
    fetching instructions/data before they are needed.

→ Drawback; wastage of memory.

→ Memory bandwidth

TRADEOFFS OF Multithreading:

→ shared memory

→ smaller cache residency of each thread.

Interconnection Networks for Parallel Computers:

→ carry data between processors & to memory.

→ Interconnects are made up of classified as static or dynamic.

(<sup>†</sup> point to point      <sup>†</sup> single path) communication

DMA.

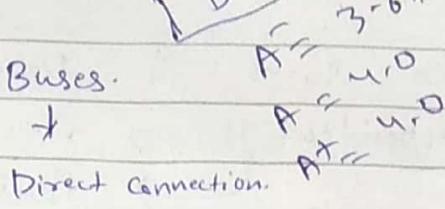
DMA → Allocate Memory.

Network Topologies : Buses.  
Buses.

$$A = 4.0 \\ B^+ = 3.33 \\ B^- = 3.67$$

$$B^- = 3.33$$

$$A^- = 3.67$$



$$B^+ = 3.33 \\ B^- = 3 \\ A^- = 3.67 \\ B^- = 2.67 \\ C^- = 2.7$$

Crossbar:

Many processor communicate each other.

Processor  $\rightarrow$  Memory  $\neq$  Connection & Communication.

Motherboard

Interconnectivity (Components of a PC)

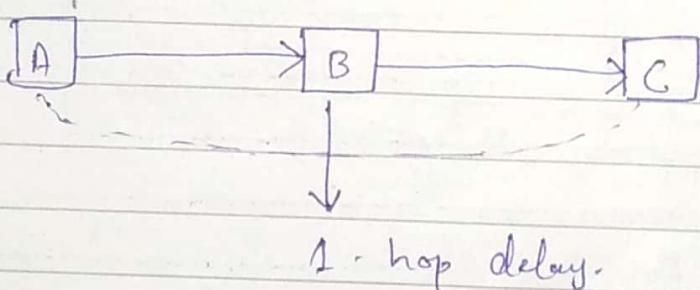
Inside a computer system

Not In Parallel

Message Passing Cost in Parallel Computers:

- $\rightarrow$  Startup Time
- $\rightarrow$  Per-hop Time (Network institution)
- $\rightarrow$  Per-Word Transfer (length of message, error & correction checking)

$\rightarrow$  B/w two system (message passing)



Babu  
Karan

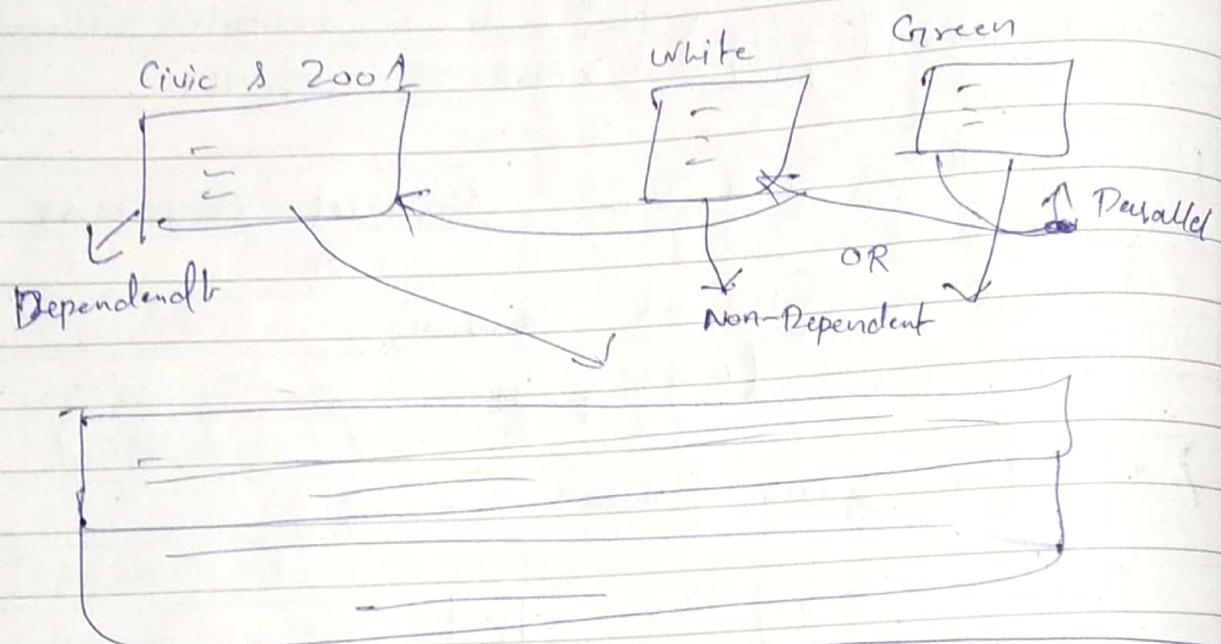
⇒ Decomposition Tasks & Dependency Graphs:  
 Task dependent ⇒ Can't be parallel.  
 ↳ Pipeline concept.

Problem decompose to tasks.

May or may not be of same size.

↳ ex: Vector Dot Product. (No Dependency).

ex: Database searching some result  
 ↳ Parallel & No Dependency.



Granularity:

Fine ← Communication Intensive.

Coarse ← Computation Intensive.

Fine-Grained Decomposition → More communication.

Coarse-Grained Decomposition → More computation

↳ Matrix multiplication.

⇒ Degree of Concurrency:

No. of tasks executing (going) parallel.  
e.g.: 4, 5.

AND = Dependent / Not Parallel

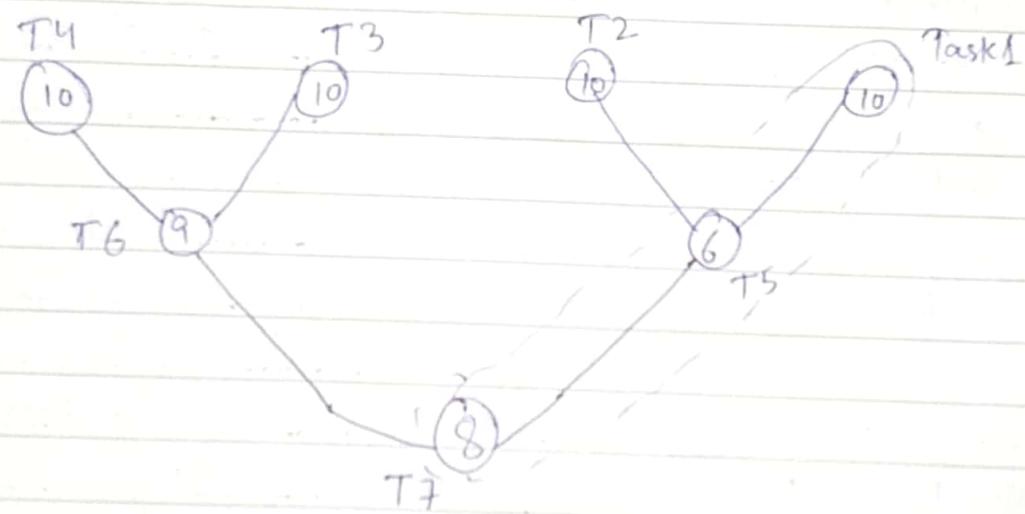
OR = Parallel.

Task Decomposition

Efficient Parallelism

Critical Path length:

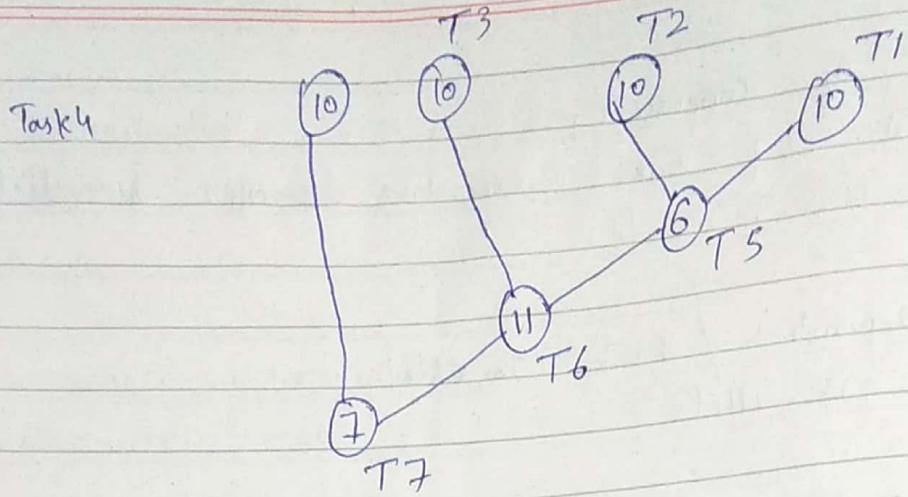
→ Length of longest path in critical task dependency.



$$\begin{aligned} \text{Total amount of work} &= 10 + 10 + 10 + 10 + 9 + 6 + 8 \\ &= 63 \end{aligned}$$

$$\text{critical path length} = 27$$

$$\text{Avg. degree of concurrency} = \frac{63}{27} = 2.33$$



Total amount of work = 64

C.P.L = 34

Avg. degree of concurrency =  $64/34 = 1.88$

### Task Interaction Graphs: (Data dependence)

→ Subtasks generally exchange data with others in a decomposition.

→ Represent data dependency to

AND in query DBMS is controlled dependency.

### Task Interaction

Control Dependency: Pipeline = Branch & Jump

ADD R<sub>1</sub>, R<sub>1</sub>, R<sub>2</sub>

BEQ R<sub>1</sub>, R<sub>3</sub>, Label

ADD R<sub>2</sub>, R<sub>3</sub>, R<sub>4</sub>

ADD has control dependency on branch

Data / Task Dependency: Parallel = Data Representation.

↳ Value doesn't exist

↳ Data dependency.

until operation is finished.

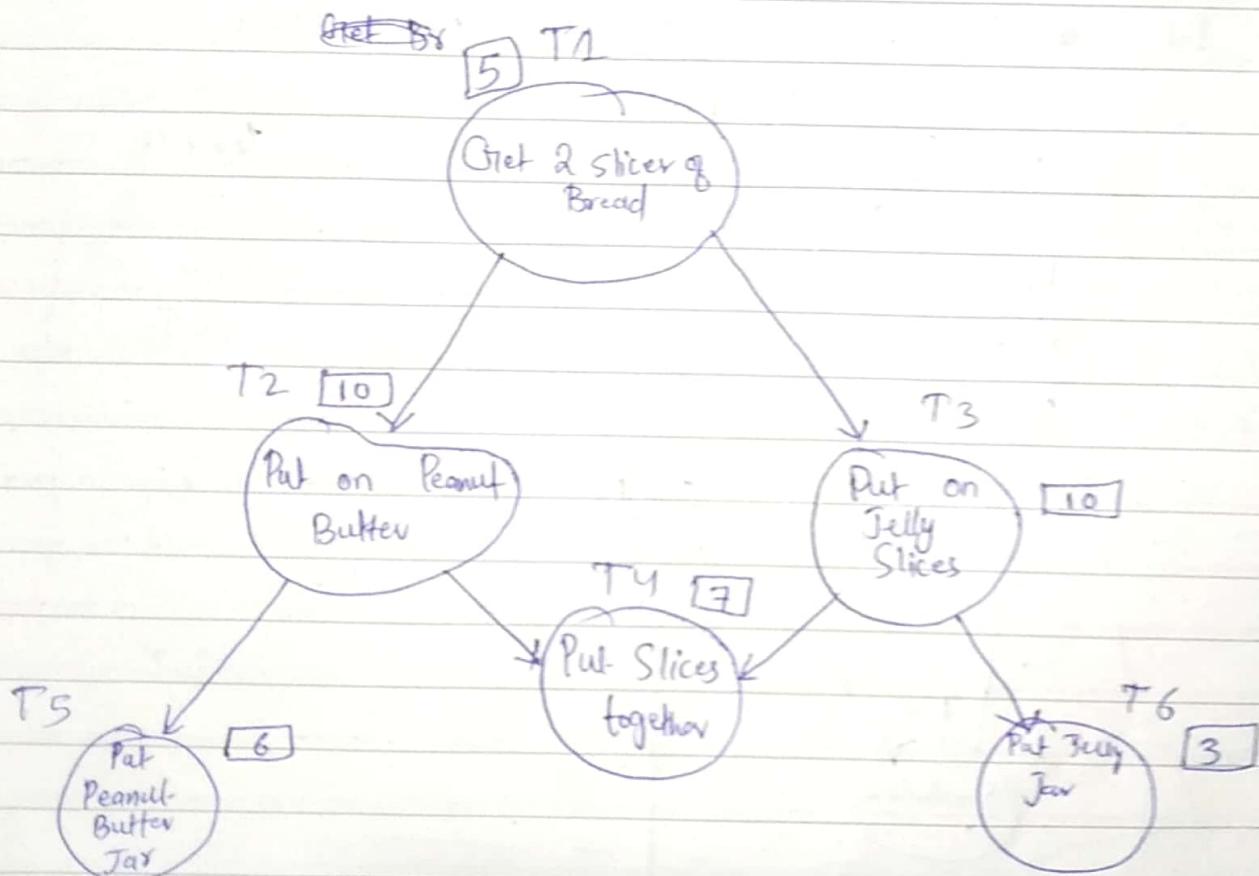
ADD R<sub>1</sub>, R<sub>2</sub>, R<sub>3</sub>

SUB R<sub>2</sub>, (R<sub>1</sub>), R<sub>8</sub>

## Task Dependency Graph

Steps for making butter and jelly sandwich.  
You must first get the bread before spreading the peanut butter jelly on the bread.  
You must first get spread peanut butter before you put away peanut butter and likewise spread on jelly before you put away jelly jar.  
And you need to spread or both peanut butter jelly before putting two slices of bread together.  
It doesn't matter if you spread peanut butter first or jelly first.

Control



Task Dependency = Data dep.

Control Dependency =

Task Interaction = Data dependency  
Task Dependency = Control dependency / dependencies.

Interactive task = 1, 3, 5

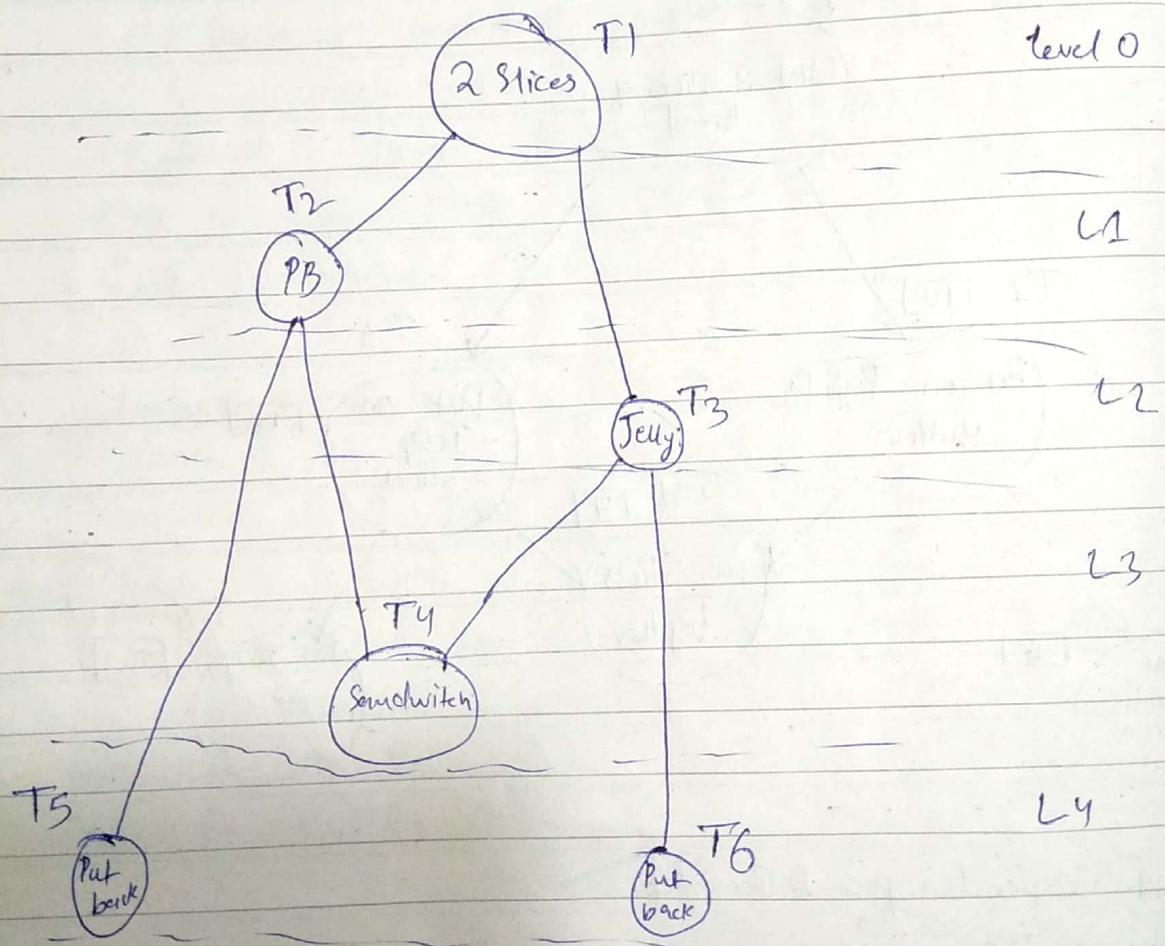
Task Dependency = 2, 3, 4, 5, 6

Total Amount of time = 41

CPL = 22

Degree of Concurrency =  $\frac{41}{22} = 1.86$ .

g) Put peanut butter first and then jelly.



## Process & Mapping

- Mappings are determined by both the task dependency.
- Minimum communication (Cause grand  $\rightarrow$  computation)
- Critical Path to process as soon as they become available.

Minimizing interaction = Coarse grain.

Dense Interaction = Fine grain

## Mapping:

- ① Mapping independent task to different processes.
- ② Assign C.P to process as soon as they become available.
- ③ Minimizing interaction b/w processes by mapping with

## Decomposition Techniques:

- ① Recursive Decomposition (Divide & Conquer)
  - ② Data Decomposition (partition data across various tasks)
    - ↳ Matrix Dot Product (MM), DB Transactions.
    - ↳ E.g.: Web searching.
- ↓
- ↳ Frequency itemset Mining
- ↳ Apriori Algorithm.

## Programming

Open MP = shared memory space parallel programming  
 threads = shared  
 multithreading = shared memory

#pragma omp parallel [clause list]  
 # pragma omp directive [clause list]

Degree of Concurrency:

num-threads( int )  
 ↳ no. of threads created.

Barrier = Synchronization = pthread-join .

default (private)

nowait =

default:

first private , private , none, last private.  
 ↑  
 initialize  
 to initial  
 value.

⇒ Synchronization:

Critical

↳ race condition

↳ Data shared among all processors

↳ Critical section one ek processor

ek waqt he data access karega.

# Shared Memory, Distributed Memory, Granularity

ordered:

The structure block is executed in the order = sequential.

Barrier:

Each thread waits until all of the other threads of a team have reached this point.

No wait:

Wait nahi karega jo barvi peek threads execute ho rhe hain. Barvi attaches a wait nahi kaega.

Schedule (type, chunk)

① Static

↓  
Har thread ke  
equally divide de  
dyo chunks me

② Dynamic

Kuch threads  
apne calculate kye.  
Agla chunks kechain  
takey me usse  
Kam kar sakon

③ Guided

Large chunk of contiguous  
iterations are allocated to  
each thread at a time.  
Khud ke compute kry GPU  
kaise dena hai task.

Sections:

Task parallel me execute hange dor for loop  
lohe nahi (iterations).

#pragma omp parallel

{ #pragma omp sections

{ #pragma omp sections  
{ task A(); }

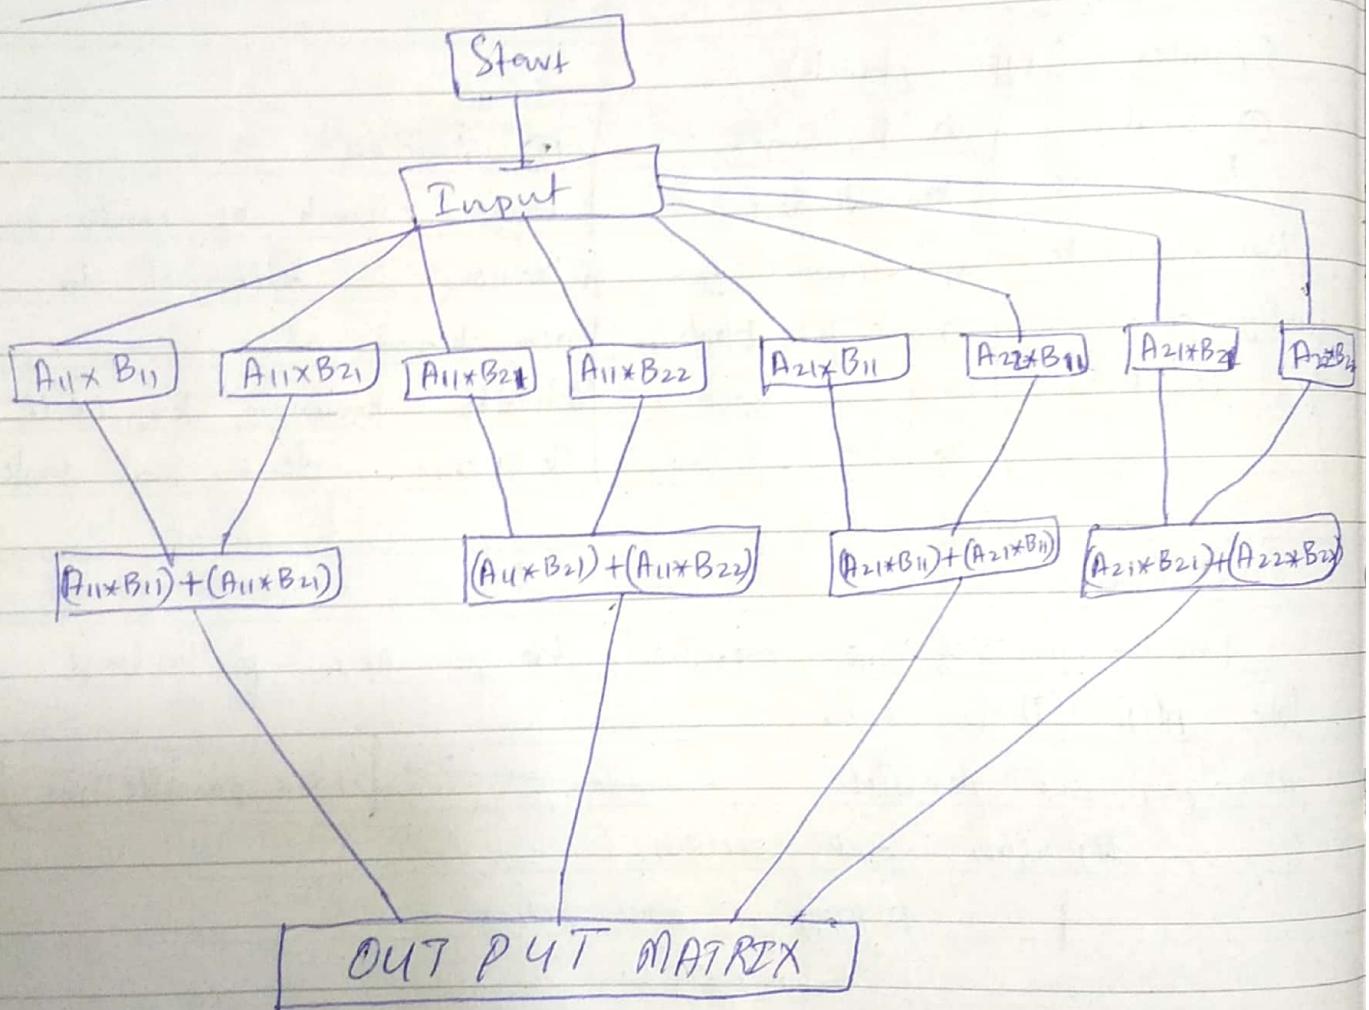
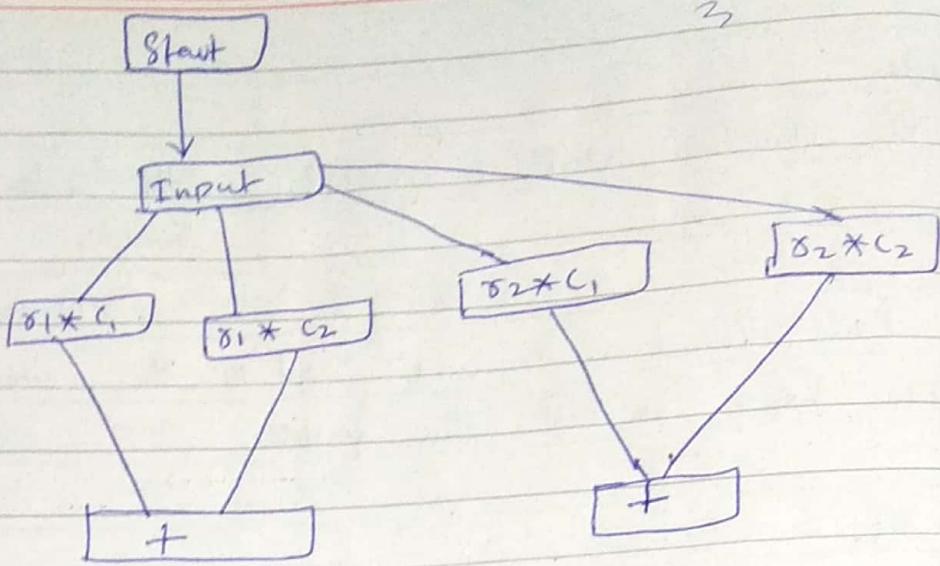
#pragma omp sections

{ task B(); }

#pragma omp sections

{ task C(); }

Task parallelism



D[0][0]  
0 1

```
main() { B[2][2] = {0};  
         A[2][2] = {0};  
         C[2][2] = {0};  
         int mult(x, y) // Function to multiply  
         { return x * y } }
```

```
int sum(x, y) // Function to sum,  
{ return x + y }
```

```
# pragma omp section  
# pragma omp parallel num-thread(8)
```

```
for(i=0; i<4; i++)  
{ # pragma omp too nested-loop  
    for(j=0; j<8; j++)  
    { C[i][j] = multiply(A[i][j], B[i][j]); }  
    D[i][j] = sum(C[i][j], C[i+1][j+1]); }
```

```
# pragma omp parallel for
```

```
for(k=0; k<8; k++)
```

DEK Sum(C[k], C[k+1])

foo

# Multi programming.

Symmetric

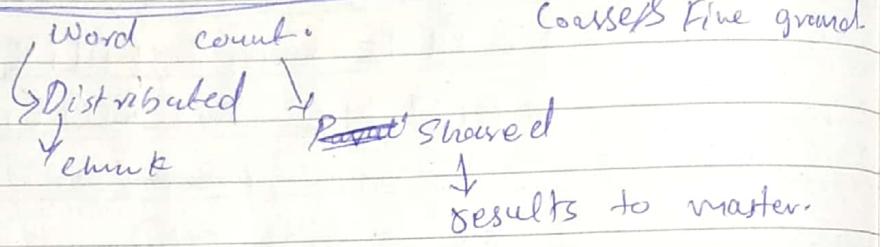
$$m \times m = [128 \times 128]$$

SMP  $\rightarrow$  Shared Memory = Coarse grain  
ASMP  $\rightarrow$  Distributed ~~Memory~~ = Fine grain

UMA:

shared array unified:  
 $\hookrightarrow$  sabhko esather mil jaxega.

Hybrid Distributed-Shared Memory:



CH #  
Instructions Types

- ① Serial
- ② Pipeline
- ③ Parallel
- ④ Jump
- ⑤ Branch
- ⑥ Iterative

SUPERSCALAR EXECUTION:

Empty	SUPERSCALAR EXECUTION:	
Partial	① Horizontal wastage.	Not fully parallel utilized
Full	② Vertical wastage.	No CPU utilize

## Data Dependency.

### DECOMPOSITIONS:

- ① Recursive Decomposition      { Divide & Conquer}
- ② Data Decomposition      { Matrix Multiplications }
  - ↳ Transitions
  - ↳ Item set example.

### ③ Exploratory:

Game next stage/step prediction.  
e.g.: TIC TAC TOE game.

- ④ Speculative Decomposition:
  - getup      { In order }
  - Get ready      | Roll Back
  - drive to work      | Switch dependent instructions
  - work      | execute some or
  - eat lunch      | all cases ~~before~~ in advance

### Theory

Scenario → Graph → Code → Open MP

Task Dependency Graph.

Interception Graph.

Symmetric	vs.	Not Asymmetric	Multiprogressing:
↓ Shared memory			Not shared memory.
↳ bus based access to common physical memory.			Master $\rightarrow P_1$
↳ In parallel computing it describes a model where parallel tasks have all same picture of memory.			Slave $\rightarrow P_3$
↳ race condition			Slave $\rightarrow P_2$
↳ semaphore			Distributed Memory:
			No CC.
			Communication network
			Each processor has its own memory.
		NUMA	times / words count ↳ memory multiplications Downloading a file

### Distributed Memory:

In hardware:

Refers to a network based memory access for physical memory that is not common.

In Parallel Programming:

Tasks can only logically "see" local machine memory and must use communications to access memory on other machines where other tasks are executing.

### Synchronization:

In simple words synchronization is waiting for one or more tasks in an application to reach the same or logically equivalent point.

\* Is a overhead / Delays.

### Granularity:

Qualitative measure of the ratio of computation to communication.

UMA = Unified Memory Access.

### → Coarse Grain:

Computational extensive.

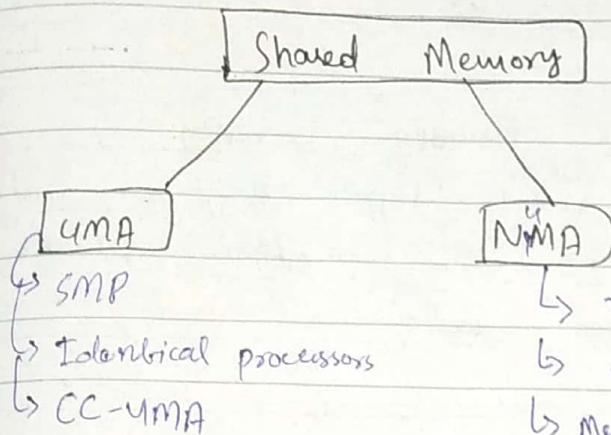
### Scalability:

Parallel system's hardware or software ability to demonstrate proportionate increase in parallel speedup with the addition of more processors.

→ Application algorithm.

→ Parallel overhead.

→ Synchronization.



↳ SMP

↳ Identical processors

↳ CC-UMA

↳ Two or More SMP

↳ Can access another SMP's memory

↳ Memory access across link is slower

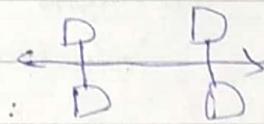
CC = Cache Coherence

Not equal access time

Equal Access Time

Shared <sup>to</sup> memory

### Hybrid Distributed-Shared Memory:



→ Cache coherent Shared Memory Processing.

another.

→ Network communications. to move data from one SMP to

→ Downloading a file in chunks & combining them.

→ Words Count

→ Matrix Multiplication.

### Threads:

Model of parallel programming, a single process can have multiple, concurrent execution paths.

→ Subroutines in main program can be executed

→ Threads communicate with each other through global memory.

OpenMP

Latency = Delay

Bandwidth = Memory access time / sec

Caches = low latency high-bandwidth.

Fetch = latency = 200 ns.

Approaches for Hiding Memory Latency:

- ① Prefetching → fetching <sup>data,</sup> results before they are actually needed.
- ② Multithreading → dividing task into multiple subtasks.
- ③ Cache to load faster.

#pragma omp parallel default(private) shared(a,b,sum)  
reduction(+:sum)

join all threads each copy of sum to parent.

Master thread.

# pragma omp for

① ordered = executed in order, sequential loop

② Barrier = Synchronization

↳ All threads wait until

Each thread waits until all threads reach at common or logically equivalent point.

③ nowait = threads completing assigned work can proceed without waiting for all threads in the team to finish.

Matrix

Multiplication:

thrust.c

{

#pragma omp parallel num\_threads(8)

#pragma omp parallel default(private)  
shared(a, b, c, dim) num\_threads(4)

#pragma omp for schedule(static)

for(i=0; i<dim; i++)  
    { assign iteration to threads.  
        for(j=0; j<dim; j++)  
            ↳ used with loops only.

    { c[i][j]=0;

        for(k=0; k<dim; k++)

            c[i][j] += a[i][k] \* b[k][j]; } }

}

①

$$a[2][2] = \begin{bmatrix} 1 & 2 \\ 0 & 1 \end{bmatrix}$$

$$b[2][2] = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}$$

i=0

j=0

$$c[0][0] = 0$$

$$\begin{array}{r} 0+2 \\ 0+1 \\ \hline 0+1 \end{array} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

K=0

$$c[0][0] += a[0][0] * b[0][0]$$

$$// 0+ = 1*0 = 0$$

$$c[0][0] += a[0][1] * b[1][0]$$

$$// 0+ = 2*1 = 2$$

$$c[0][0] += a[0][0] * b[0][1]$$

$$// 0+ 2*1 = 2$$

$$c[0][1] += a[0][1] * b[1][1]$$

$$// 1+ 2 = 3$$

$$c[0][1]$$

$$c[1][1]$$